



Mixed-format diagnostic: both query families trained; direct answer sees no gold trace

Agreement with actual answer gradient



Parameter-gradient magnitude



Does actual descent improve true rules?

