

Mixed-format diagnostic: both query families trained; direct answer sees no gold trace

Agreement with actual answer gradient Parameter-gradient magnitude Does actual descent improve true rules?

